

Skills Network Labs

File Edit View Run Kernel Git Tabs Settings Help

Launcher 8.4.1_MomentumwithPoly... 8.3.1.initializationsame.ipynb Python



Momentum

Objective for this Notebook

1. Learn Saddle Points, Local Minima, and Noise

Table of Contents

In this lab, you will deal with several problems associated with optimization and see how momentum can improve your results.

- Saddle Points
- Local Minima
- Noise

Estimated Time Needed: 25 min

Preparation

Import the following libraries that you'll use for this lab:

```
[ ]: # These are the Libraries that will be used for this Lab.

import torch_
import torch.nn as nn
import matplotlib.pyplot as plt
import numpy as np

torch.manual_seed(0)
***
```

This function will plot a cubic function and the parameter values obtained via Gradient Descent.

```
[ ]: # Plot the cubic

def plot_cubic(w, optimizer):
    LOSS = []
    # parameter values...
    W = torch.arange(-4, 4, 0.1)
    # plot the loss function.
    for w.state_dict()['linear.weight'][0] in W:
        LOSS.append(cubic(w(torch.tensor([[1.0]]))).item())
    w.state_dict()['linear.weight'][0] = 4.0
    n_epochs = 10
    parameter = []
    loss_list = []

    # n_epochs
    # Use PyTorch custom module to implement a polynomial function
    for n in range(n_epochs):
        optimizer.zero_grad_()
        loss = cubic(w(torch.tensor([[1.0]])))
        loss_list.append(loss)
        parameter.append(w.state_dict()['linear.weight'][0].detach().data.item())
        loss.backward()
        optimizer.step()

    plt.plot(parameter, loss_list, 'ro', label='parameter values')
    plt.plot(W.numpy(), LOSS, label='objective function')
    plt.xlabel('w')
    plt.ylabel('l(w)')
    plt.legend()
```

This function will plot a 4th order function and the parameter values obtained via Gradient Descent. You can also add Gaussian noise with a standard deviation determined by the parameter std .

```
[ ]: # Plot the fourth order function and the parameter values

def plot_fourth_order(w, optimizer, std=0, color='r', paramlabel='parameter values', objfun=True):
    W = torch.arange(-4, 6, 0.1)
    LOSS = []
    for w.state_dict()['linear.weight'][0] in W:
        LOSS.append(fourth_order(w(torch.tensor([[1.0]]))).item())
    w.state_dict()['linear.weight'][0] = 6
    n_epochs = 100
    parameter = []
    loss_list = []

    #n_epochs
    for n in range(n_epochs):
        optimizer.zero_grad_()
        loss = fourth_order(w(torch.tensor([[1.0]]))) + std * torch.randn(1, 1)
```

```

        loss_list.append(loss)
        parameter.append(w.state_dict()['linear.weight'][0].detach().data.item())
        optimizer.step()

        # Plotting
        if objfun:
            plt.plot(W.numpy(), LOSS, label='objective function')
            plt.plot(parameter, loss_list, 'ro', label=paramlabel, color=color)
            plt.xlabel('w')
            plt.ylabel('l(w)')
            plt.legend()

```

This is a custom module. It will behave like a single parameter value. We do it this way so we can use PyTorch's build-in optimizers .

```
[ ]: # Create a Linear model
class one_param(nn.Module):
    # Constructor
    def __init__(self, input_size, output_size):
        super(one_param, self).__init__()
        self.linear = nn.Linear(input_size, output_size, bias=False)

    # Prediction
    def forward(self, x):
        yhat = self.linear(x)
        return yhat
```

We create an object `w`, when we call the object with an input of one, it will behave like an individual parameter value. i.e `w(1)` is analogous to `w`

```
[ ]: # Create a one_param object
w = one_param(1, 1)
```

Saddle Points

Let's create a cubic function with Saddle points

```
[ ]: # Define a function to output a cubic..
def cubic(yhat):
    out = yhat ** 3
    return out
```

We create an optimizer with no momentum term

```
[ ]: # Create a optimizer without momentum
optimizer = torch.optim.SGD(w.parameters(), lr=0.01, momentum=0)
```

We run several iterations of stochastic gradient descent and plot the results. We see the parameter values get stuck in the saddle point.

```
[ ]: # Plot the model
plot_cubic(w, optimizer)
```

we create an optimizer with momentum term of 0.9

```
[ ]: # Create a optimizer with momentum
optimizer = torch.optim.SGD(w.parameters(), lr=0.01, momentum=0.9)
```

We run several iterations of stochastic gradient descent with momentum and plot the results. We see the parameter values do not get stuck in the saddle point.

```
[ ]: # Plot the model
plot_cubic(w, optimizer)
```

Local Minima

Did you know? IBM Watson Studio lets you build and deploy an AI solution, using the best of open source and IBM software and giving your team a single environment to work in. [Learn more here](#).

In this section, we will create a fourth order polynomial with a local minimum at 4 and a global minimum a -2. We will then see how the momentum parameter affects convergence to a global minimum. The fourth order polynomial is given by:

```
[ ]: # Create a function to calculate the fourth order polynomial..
def fourth_order(yhat):
    out = torch.mean(2 * (yhat ** 4) - 9 * (yhat ** 3) - 21 * (yhat ** 2) + 88 * yhat + 48)
    return out
```

We create an optimizer with no momentum term. We run several iterations of stochastic gradient descent and plot the results. We see the parameter values get stuck in the local minimum.

```
[ ]: # Make the prediction without momentum
optimizer = torch.optim.SGD(w.parameters(), lr=0.001)
plot_fourth_order(w, optimizer)
```

We create an optimizer with a momentum term of 0.9. We run several iterations of stochastic gradient descent and plot the results. We see the parameter values reach a global minimum.

```
[ ]: # Make the prediction with momentum
optimizer = torch.optim.SGD(w.parameters(), lr=0.001, momentum=0.9)
plot_fourth_order(w, optimizer)
```

Noise

In this section, we will create a fourth order polynomial with a local minimum at 4 and a global minimum a -2, but we will add noise to the function when the Gradient is calculated. We will then see how the momentum parameter affects convergence to a global minimum.

with no momentum, we get stuck in a local minimum

```
[ ]: # Make the prediction without momentum when there is noise  
optimizer = torch.optim.SGD(w.parameters(), lr=0.001)  
plot_fourth_order(w, optimizer, std=10)
```

with momentum, we get to the global minimum

```
[ ]: # Make the prediction with momentum when there is noise  
optimizer = torch.optim.SGD(w.parameters(), lr=0.001, momentum=0.9)  
plot_fourth_order(w, optimizer, std=10)
```

Practice

Create two SGD objects with a learning rate of 0.001. Use the default momentum parameter value for one and a value of 0.9 for the second. Use the function `plot_fourth_order` with an `std=100`, to plot the different steps of each. Make sure you run the function on two independent cells.

```
[ ]: # Practice: Create two SGD optimizer with Lr = 0.001, and one without momentum and the other with momentum = 0.9. Plot the result out.  
# Type your code here
```

Double-click here for the solution.



About the Authors:

Joseph Santarcangelo has a PhD in Electrical Engineering, his research focused on using machine learning, signal processing, and computer vision to determine how videos impact human cognition. Joseph has been working for IBM since he completed his PhD.

Other contributors: [Michelle Carey](#), [Mavis Zhou](#)

Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-09-23	2.0	Srishti	Migrated Lab to Markdown and added to course repo in GitLab

© IBM Corporation 2020. All rights reserved.

